

```
In [133... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import keras
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from keras.datasets import mnist
from keras.models import Sequential
from sklearn.model_selection import GridSearchCV
from sklearn.neural_network import MLPClassifier
from keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D,
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
from keras.optimizers import Adam
```

```
In [134... df = pd.read_csv(r'C:\Users\ayaYM\Downloads\heart-2.csv')
```

```
In [135... print(df.head())
print(df.info())
print(df.describe())
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1025 entries, 0 to 1024
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	age	1025 non-null	int64
1	sex	1025 non-null	int64
2	cp	1025 non-null	int64
3	trestbps	1025 non-null	int64
4	chol	1025 non-null	int64
5	fbs	1025 non-null	int64
6	restecg	1025 non-null	int64
7	thalach	1025 non-null	int64
8	exang	1025 non-null	int64
9	oldpeak	1025 non-null	float64
10	slope	1025 non-null	int64
11	ca	1025 non-null	int64
12	thal	1025 non-null	int64
13	target	1025 non-null	int64

```
dtypes: float64(1), int64(13)
```

```
memory usage: 112.2 KB
```

```
None
```

	age	sex	cp	trestbps	chol	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	54.434146	0.695610	0.942439	131.611707	246.000000	
std	9.072290	0.460373	1.029641	17.516718	51.59251	
min	29.000000	0.000000	0.000000	94.000000	126.00000	
25%	48.000000	0.000000	0.000000	120.000000	211.00000	
50%	56.000000	1.000000	1.000000	130.000000	240.00000	
75%	61.000000	1.000000	2.000000	140.000000	275.00000	
max	77.000000	1.000000	3.000000	200.000000	564.00000	

	fbs	restecg	thalach	exang	oldpeak	\
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	
mean	0.149268	0.529756	149.114146	0.336585	1.071512	
std	0.356527	0.527878	23.005724	0.472772	1.175053	
min	0.000000	0.000000	71.000000	0.000000	0.000000	
25%	0.000000	0.000000	132.000000	0.000000	0.000000	
50%	0.000000	1.000000	152.000000	0.000000	0.800000	
75%	0.000000	1.000000	166.000000	1.000000	1.800000	
max	1.000000	2.000000	202.000000	1.000000	6.200000	

	slope	ca	thal	target
--	-------	----	------	--------

count	1025.000000	1025.000000	1025.000000	1025.000000
mean	1.385366	0.754146	2.323902	0.513171
std	0.617755	1.030798	0.620660	0.500070
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	2.000000	0.000000
50%	1.000000	0.000000	2.000000	1.000000
75%	2.000000	1.000000	3.000000	1.000000
max	2.000000	4.000000	3.000000	1.000000

In [136... `print(df.shape)`

(1025, 14)

In [137... `df.head(15)`

Out[137...

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
<b>0</b>	52	1	0	125	212	0	1	168	0	1.0	2	2	3
<b>1</b>	53	1	0	140	203	1	0	155	1	3.1	0	0	3
<b>2</b>	70	1	0	145	174	0	1	125	1	2.6	0	0	3
<b>3</b>	61	1	0	148	203	0	1	161	0	0.0	2	1	3
<b>4</b>	62	0	0	138	294	1	1	106	0	1.9	1	3	2
<b>5</b>	58	0	0	100	248	0	0	122	0	1.0	1	0	2
<b>6</b>	58	1	0	114	318	0	2	140	0	4.4	0	3	1
<b>7</b>	55	1	0	160	289	0	0	145	1	0.8	1	1	3
<b>8</b>	46	1	0	120	249	0	0	144	0	0.8	2	0	3
<b>9</b>	54	1	0	122	286	0	0	116	1	3.2	1	2	2
<b>10</b>	71	0	0	112	149	0	1	125	0	1.6	1	0	2
<b>11</b>	43	0	0	132	341	1	0	136	1	3.0	1	0	3
<b>12</b>	34	0	1	118	210	0	1	192	0	0.7	2	0	2
<b>13</b>	51	1	0	140	298	0	1	122	1	4.2	1	3	3
<b>14</b>	52	1	0	128	204	1	1	156	1	1.0	1	0	0

In [138... `df.describe()`

Out[138...

	age	sex	cp	trestbps	chol	fbs	rt
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.0
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	0.5
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	0.5
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.0
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	0.0
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.0
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	1.0
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.0

In [139...


```
df.isnull().sum()
```

Out[139...

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

In [140...

```
plt.figure(figsize=(10, 6))
sns.heatmap(df.select_dtypes(include=['number']).corr(),
            annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Matrix (Numeric Columns)")
plt.show()
```



```
In [141... X = df.drop('target', axis=1) # Features
            y = df['target']      # Target

            X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
```

```
In [142... scaler = StandardScaler()
            X_train_scaled = scaler.fit_transform(X_train)
            X_test_scaled = scaler.transform(X_test)
```

```
In [143... print(df['target'].value_counts())
```


```
target
1    526
0    499
Name: count, dtype: int64
```


```
In [144... df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1025 non-null   int64   
 1   sex         1025 non-null   int64   
 2   cp          1025 non-null   int64   
 3   trestbps    1025 non-null   int64   
 4   chol        1025 non-null   int64   
 5   fbs         1025 non-null   int64   
 6   restecg     1025 non-null   int64   
 7   thalach     1025 non-null   int64   
 8   exang       1025 non-null   int64   
 9   oldpeak     1025 non-null   float64  
10  slope       1025 non-null   int64   
11  ca          1025 non-null   int64   
12  thal        1025 non-null   int64   
13  target      1025 non-null   int64   
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
In [145... sns.countplot(x='target', data=df)
plt.show()

# Example boxplot for age vs target
sns.boxplot(x='target', y='age', data=df)
plt.show()
```





### SVM models

In [146...

```
svm_models = {
    "SVC (RBF kernel)": SVC(kernel='rbf', random_state=42),
    "Linear SVM": SVC(kernel='linear', random_state=42),
    "Polynomial SVM": SVC(kernel='poly', degree=3, random_state=42)
}

# Train, predict, and evaluate each model
for name, model in svm_models.items():
    print(f"\n--- {name} ---")

    # Train
    model.fit(X_train_scaled, y_train)

    # Predict
    y_pred = model.predict(X_test_scaled)

    # Accuracy
    acc = accuracy_score(y_test, y_pred)
    print(f"Accuracy on test set: {acc:.4f}")
```

```
--- SVC (RBF kernel) ---  
Accuracy on test set: 0.9058
```

```
--- Linear SVM ---  
Accuracy on test set: 0.8474
```

```
--- Polynomial SVM ---  
Accuracy on test set: 0.9123
```

SVC (RBF kernel) → 90.6% accuracy  
The RBF (Radial Basis Function) kernel performs very well on this dataset.  
This suggests the decision boundary is nonlinear, and RBF can model it better than a straight line.

Linear SVM → 84.7% accuracy  
The linear kernel has lower performance compared to the nonlinear kernels.  
This suggests that the relationship between features and target in the heart disease dataset is not purely linear – a linear decision boundary cannot separate the classes perfectly.

Polynomial SVM → 91.2% accuracy  
The polynomial kernel gives the highest accuracy among the three models.  
This indicates that the data has some complex patterns and interactions that the polynomial kernel can capture effectively.

In [147...

```
# Confusion matrix  
cm = confusion_matrix(y_test, y_pred)  
print("Confusion Matrix:")  
print(cm)
```

Confusion Matrix:

```
[[133  17]  
 [ 10 148]]
```

133: True negatives → class 0 correctly predicted as 0

17: False positives → class 0 wrongly predicted as 1

10: False negatives → class 1 wrongly predicted as 0

148: True positives → class 1 correctly predicted as 1

polynomial SVM is performing best, followed closely by the RBF SVM, while the linear SVM is lagging.

In [148...

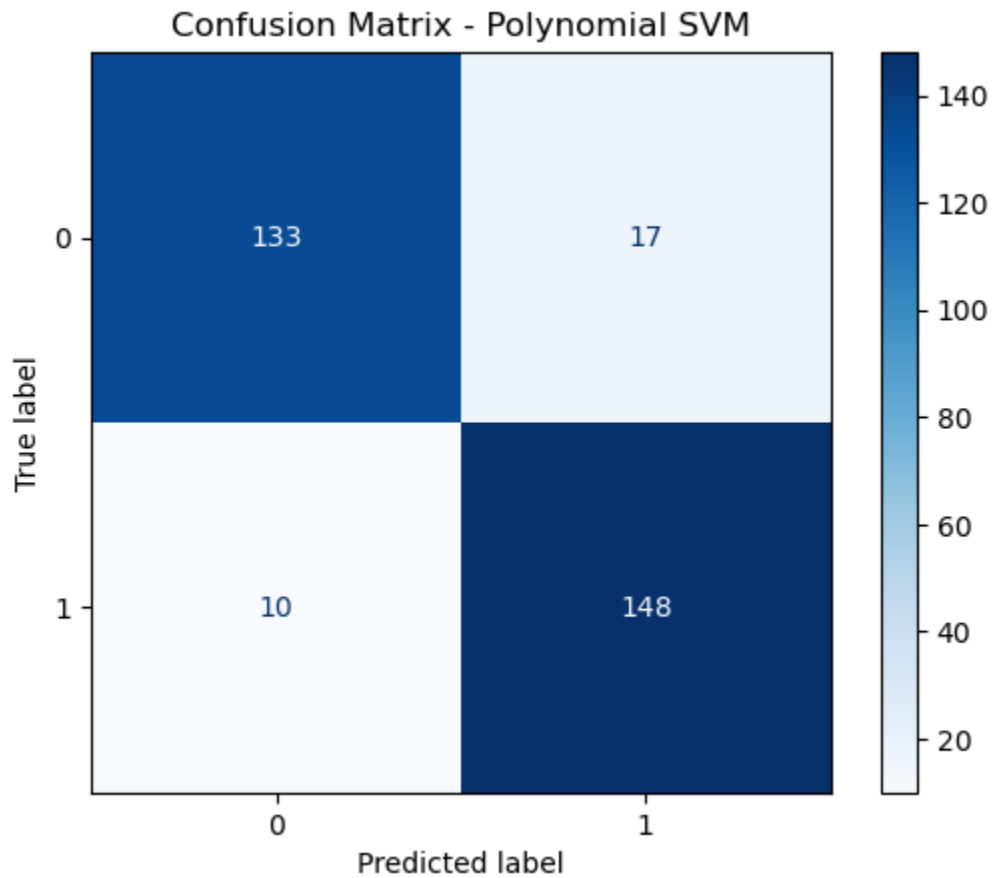
```
print("Classification Report:")  
print(classification_report(y_test, y_pred))  
  
# Plot confusion matrix  
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes)
```



```
disp.plot(cmap=plt.cm.Blues)  
plt.title(f"Confusion Matrix - {name}")  
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.89	0.91	150
1	0.90	0.94	0.92	158
accuracy			0.91	308
macro avg	0.91	0.91	0.91	308
weighted avg	0.91	0.91	0.91	308



Balanced performance on both classes

High overall accuracy (91%)

No obvious sign of extreme class imbalance or favoring one class

Neural Network Implementation Using Keras

GridSearchCV for best ALPHA and LEARNING RATE

In [149...

```
param_grid = {  
    'alpha': [0.0001, 0.001, 0.01],  
    'learning_rate_init' : [0.001, 0.01, 0.1]
```

```
}
grid_search = GridSearchCV(MLPClassifier(max_iter=1000, random_state=42), param_grid)
grid_search.fit(X_train_scaled, y_train)
alpha_best = grid_search.best_params_['alpha']
learning_rate_best = grid_search.best_params_['learning_rate_init']
print(f"\nBest Alpha from GridSearchCV: {alpha_best}")
print(f"\nBest Learning_rate from GridSearchCV: {learning_rate_best}")
```

Fitting 5 folds for each of 9 candidates, totalling 45 fits

Best Alpha from GridSearchCV: 0.0001

Best Learning\_rate from GridSearchCV: 0.001

```
In [150... model = Sequential([
    Dense(64, input_shape=(13,), activation='relu'), #input layer
    Dense(32, activation='relu'),                    #hidden layer
    Dense(2, activation='softmax')                    #output layer
])
```

c:\Users\ayaYM\anaconda3\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [151... model.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 64)	896
dense_25 (Dense)	(None, 32)	2,080
dense_26 (Dense)	(None, 2)	66

Total params: 3,042 (11.88 KB)

Trainable params: 3,042 (11.88 KB)

Non-trainable params: 0 (0.00 B)

preparing the model for training :

which optimizer to use

which loss function to minimize

which evaluation metrics to track during training & validation

```
In [152... model.compile(Adam(learning_rate=learning_rate_best), loss='sparse_categorical_crossentropy')
```

Training the model

```
In [153... history = model.fit(X_train_scaled, y_train, batch_size=10, epochs=20, shuffle=True
```

Epoch 1/20  
72/72 - 5s - 69ms/step - accuracy: 0.7545 - loss: 0.5214 - val\_accuracy: 0.8279 - val\_loss: 0.3752  
Epoch 2/20  
72/72 - 0s - 6ms/step - accuracy: 0.8550 - loss: 0.3506 - val\_accuracy: 0.8571 - val\_loss: 0.3127  
Epoch 3/20  
72/72 - 0s - 5ms/step - accuracy: 0.8745 - loss: 0.3003 - val\_accuracy: 0.8701 - val\_loss: 0.2902  
Epoch 4/20  
72/72 - 0s - 5ms/step - accuracy: 0.9010 - loss: 0.2691 - val\_accuracy: 0.8701 - val\_loss: 0.2751  
Epoch 5/20  
72/72 - 0s - 5ms/step - accuracy: 0.9079 - loss: 0.2399 - val\_accuracy: 0.8961 - val\_loss: 0.2604  
Epoch 6/20  
72/72 - 0s - 4ms/step - accuracy: 0.9191 - loss: 0.2151 - val\_accuracy: 0.9091 - val\_loss: 0.2436  
Epoch 7/20  
72/72 - 0s - 5ms/step - accuracy: 0.9372 - loss: 0.1926 - val\_accuracy: 0.9123 - val\_loss: 0.2343  
Epoch 8/20  
72/72 - 0s - 4ms/step - accuracy: 0.9470 - loss: 0.1671 - val\_accuracy: 0.9286 - val\_loss: 0.2163  
Epoch 9/20  
72/72 - 0s - 4ms/step - accuracy: 0.9512 - loss: 0.1515 - val\_accuracy: 0.9448 - val\_loss: 0.2114  
Epoch 10/20  
72/72 - 0s - 5ms/step - accuracy: 0.9637 - loss: 0.1285 - val\_accuracy: 0.9513 - val\_loss: 0.1955  
Epoch 11/20  
72/72 - 0s - 5ms/step - accuracy: 0.9693 - loss: 0.1109 - val\_accuracy: 0.9545 - val\_loss: 0.1813  
Epoch 12/20  
72/72 - 0s - 5ms/step - accuracy: 0.9763 - loss: 0.0998 - val\_accuracy: 0.9578 - val\_loss: 0.1680  
Epoch 13/20  
72/72 - 0s - 5ms/step - accuracy: 0.9833 - loss: 0.0917 - val\_accuracy: 0.9643 - val\_loss: 0.1576  
Epoch 14/20  
72/72 - 0s - 4ms/step - accuracy: 0.9847 - loss: 0.0752 - val\_accuracy: 0.9610 - val\_loss: 0.1489  
Epoch 15/20  
72/72 - 0s - 4ms/step - accuracy: 0.9847 - loss: 0.0689 - val\_accuracy: 0.9643 - val\_loss: 0.1391  
Epoch 16/20  
72/72 - 0s - 5ms/step - accuracy: 0.9902 - loss: 0.0585 - val\_accuracy: 0.9643 - val\_loss: 0.1304  
Epoch 17/20  
72/72 - 0s - 4ms/step - accuracy: 0.9930 - loss: 0.0514 - val\_accuracy: 0.9643 - val\_loss: 0.1284  
Epoch 18/20  
72/72 - 0s - 5ms/step - accuracy: 0.9930 - loss: 0.0445 - val\_accuracy: 0.9708 - val\_loss: 0.1164  
Epoch 19/20  
72/72 - 0s - 4ms/step - accuracy: 0.9902 - loss: 0.0405 - val\_accuracy: 0.9773 - val

```
_loss: 0.1159
Epoch 20/20
72/72 - 0s - 4ms/step - accuracy: 0.9944 - loss: 0.0360 - val_accuracy: 0.9773 - val
_loss: 0.1129
```

```
In [154... test_loss, test_accuracy = model.evaluate(X_test_scaled, y_test, verbose=0)
print(f'Test Loss: {test_loss:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')
```

```
Test Loss: 0.1129
Test Accuracy: 0.9773
```

```
In [155... # Get predictions
y_pred_probs = model.predict(X_test_scaled)
y_pred = y_pred_probs.argmax(axis=1)

# Classification report
print(classification_report(y_test, y_pred))

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(conf_matrix)
```

```
10/10 ————— 0s 10ms/step
```

	precision	recall	f1-score	support
0	0.97	0.98	0.98	150
1	0.98	0.97	0.98	158
accuracy			0.98	308
macro avg	0.98	0.98	0.98	308
weighted avg	0.98	0.98	0.98	308

```
Confusion Matrix:
[[147  3]
 [ 4 154]]
```

The neural network outperforms all SVM models across every metric — accuracy, precision, recall, and F1-score. It makes much fewer mistakes (7 vs 27 errors).


```
In [156... plt.figure(figsize=(12, 5))

# Loss plot
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss', color='red')
plt.plot(history.history['val_loss'], label='Validation Loss', color='orange')
plt.title('Loss over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

# Accuracy plot
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy', color='blue')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy', color='green')
```

```
plt.title('Accuracy over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



In [157...

```
nn_test_accuracy = test_accuracy # from model.evaluate


svm_accuracies = {}
for name, svm_model in svm_models.items():
    svm_model.fit(X_train_scaled, y_train)
    y_pred = svm_model.predict(X_test_scaled)
    acc = accuracy_score(y_test, y_pred)
    svm_accuracies[name] = acc

# Step 2: Combine all into one dictionary
all_accuracies = svm_accuracies.copy()
all_accuracies['Neural Network'] = nn_test_accuracy

plt.figure(figsize=(10, 6))
colors = sns.color_palette('pastel')[0:4]
bars = plt.bar(all_accuracies.keys(), all_accuracies.values(), color=colors, edgeco

# Step 4: Add accuracy labels on top (in %)
for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width() / 2,
             height + 0.01,
             f"{height * 100:.2f}%",
             ha='center',
             va='bottom',
             fontsize=11,
             fontweight='bold')
```

```
plt.ylim(0, 1.1)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.xlabel('Model', fontsize=12)
plt.ylabel('Test Accuracy', fontsize=12)
plt.title(' Model Comparison: SVM vs Neural Network', fontsize=14, fontweight='bold')
plt.xticks(rotation=15, fontsize=11)
plt.yticks(fontsize=11)
plt.tight_layout()
plt.show()
```



#### Summary of Key Findings:

Neural Network is the best-performing model with the highest accuracy (97.73%) and a relatively low test loss (0.1129). It seems to generalize well, providing excellent results.

SVM with RBF Kernel is also very strong, achieving 90.58% accuracy, but it's still slightly behind the neural network.

Polynomial SVM is competitive with the RBF kernel, yielding an accuracy of 91.23%. However, it is still a bit less effective than the RBF kernel in this particular test.

Linear SVM performs the worst among the three, with an accuracy of 84.74%. It's not well-suited for datasets with non-linear relationships, which is likely why its performance is lower.